



Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 19 of 21

In the keys to the cabinet are on the table, which word does the verb agree with?

- ☐ Cabinet, because it is the noun closest to the verb
- ☐ Table, because it is the last noun in the sentence
- ☐ Keys, the true subject, even though the nearer noun cabinet sits between them
- ☐ The, because it is the most frequent word in the sentence

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 20 of 21

Which statement about a set of attention weights for one word is correct?

- ☐ The weights are always equal for every word
- ☐ Exactly one weight is one and all the others are zero
- ☐ The weights are positive and add up to one across all the words
- ☐ The weights can be negative and need not add up to anything

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Kalvium

Ask Gemini

kalvium.community/assessment/intro_genai4_v2_mpa3/attempt/c4435664-f366-4e37-bfc7-d2e41c97ba4a

School

 Foundations of Generative AI | MPA - 3 | 2026 - 27

57m left

Finish Assessment

Your camera feed, audio and screen are being proctored.



Instructions

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Question 21 of 21

You are exploring the three ideas: how **attention** decides what each word looks at, why order needs **positional information** added back, and how a **diffusion** model turns noise into a clean signal step by step. You will build each mechanism yourself in NumPy over the small, provided data and query it inside the notebook. Nothing is trained and no model is downloaded: the point is to make the concepts concrete with your own code. Everything runs as plain Python. **No API key, no network call, and no external model is needed.**

Mandatory Environment

You **must** solve this in **Google Colab** (open colab.research.google.com, then choose New notebook). Sign in so your work saves. Run **every** cell so that all outputs and your Text cell are visible in the exported PDF. You do not need to install anything; NumPy is already available.

Cell A (Do Not Change)

Paste this into the **first cell** exactly as it is and run it. It sets up the entire prerequisite: the three word vectors for the sentence "cat chased dog", a positional-encoding matrix, a permutation that reorders the sentence to "dog chased cat", a clean target signal for the diffusion task, a seeded random generator, and a mean-squared-error helper. Do not modify anything in this cell. Cell A defines no attention, no positional demonstration, and no diffusion loop; that is what you build.

Previous

kalvium.community is sharing your screen.


Stop sharing

[Hide](#)

Next

33°C Sunny

Search



ENG IN

12:04 09-10-2026



Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

```
import numpy as np
np.set_printoptions(precision=3, suppress=True)

# Sentence "cat chased dog": one row vector per word (these are the token embeddings).
tokens = ["cat", "chased", "dog"]
X = np.array([
    [1.0, 0.2, 0.0, 0.1], # cat
    [0.1, 1.0, 0.3, 0.0], # chased
    [0.9, 0.1, 0.0, 0.2], # dog
])

# Positional encoding: one row per POSITION (0, 1, 2). Added to a word to mark where it sits.
P = np.array([
    [0.0, 0.0, 0.0, 0.0], # position 0
    [0.5, 0.4, 0.3, 0.2], # position 1
    [1.0, 0.8, 0.6, 0.4], # position 2
])

# Reordering the sentence to "dog chased cat" (same words, different order).
perm = [2, 1, 0]

# Diffusion target: the clean signal a prompt is asking the model to recover.
x_clean = np.array([0.0, 0.2, 0.4, 0.6, 0.8, 1.0, 0.8, 0.6,
                    0.4, 0.2, 0.0, -0.2, -0.4, -0.2, 0.0, 0.2])
```

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

```
# Diffusion target: the clean signal a prompt is asking the model to recover.
x_clean = np.array([0.0, 0.2, 0.4, 0.6, 0.8, 1.0, 0.8, 0.6,
                   0.4, 0.2, 0.0, -0.2, -0.4, -0.2, 0.0, 0.2])

rng = np.random.default_rng(0)

def mse(a, b):
    return float(np.mean((a - b) ** 2))

print("tokens:", tokens)
print("X shape:", X.shape, "| P shape:", P.shape, "| x_clean length:", len(x_clean))
```

Task 1 : Build Attention from Scratch

Write this cell from scratch. Define `softmax(v)` that turns a vector of scores into positive weights that sum to one, and `attention(M)` that, for a matrix of word vectors `M`, scores every word against every word (the dot products `M @ M.T`), turns each row of scores into weights with your `softmax`, and returns the new word vectors (the weighted blend `W @ M`) together with the weight grid `W`. Run it on `X` and print the weight grid and the row sums so it is clear each row sums to one.

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next

Kalvium

Ask Gemini

kalvium.community/assessment/intro_genai4_v2_mpa3/attempt/c4435664-f366-4e37-bfc7-d2e41c97ba4a

School

 Foundations of Generative AI | MPA - 3 | 2026 - 27

57m left

Finish Assessment

Your camera feed, audio and screen are being proctored.



Instructions

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Task 1 : Build Attention from Scratch

Write this cell from scratch. Define `softmax(v)` that turns a vector of scores into positive weights that sum to one, and `attention(M)` that, for a matrix of word vectors `M`, scores every word against every word (the dot products `M @ M.T`), turns each row of scores into weights with your `softmax`, and returns the new word vectors (the weighted blend `W @ M`) together with the weight grid `W`. Run it on `X` and print the weight grid and the row sums so it is clear each row sums to one.

```
# TODO 1: Define softmax(v) and attention(M) from scratch. Run attention on X and print the
#           weight grid W and its row sums (each row must sum to 1).
```

Task 2 : Read the Attention Weights

Write this cell from scratch. Using the weight grid from Task 1, for each word print which other word it attends to most strongly and the weight value. Then confirm in one printed line that every row of weights sums to one. Notice whether the most-attended word is always the nearest one.

```
# TODO 2: For each word print the word it attends to most and that weight, then print a single
#           line confirming every row of the weight grid sums to 1.
```

Previous

kalvium.community is sharing your screen.


Stop sharing

Hide

Next

33°C Sunny

Search



ENG IN 12:05 09-10-2026



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Task 3 : Show Why Positional Information Is Needed

Write this cell from scratch. First show that bare attention is **order-blind**: compute the attention output for `X` and for the reordered sentence `X[perm]`, and check that reordering the words only reorders the outputs (the output for `X[perm]` equals the output for `X` with its rows reordered by `perm`). Then add positional information by using `X + P`, and repeat the same check on `X + P` versus `X[perm] + P` (the same words but now sitting at different positions). Print the boolean result of both checks so the contrast is visible.

```
# TODO 3: Print two booleans. First: bare attention is order-blind (reordering words only
#         reorders the outputs). Second: after adding P, that no longer holds (positions matter).
```

Task 4 : Denoise a Signal Step by Step

Write this cell from scratch. Model diffusion as cleaning up noise over many small steps toward the prompt's target. Starting from pure noise `rng.normal(size=x_clean.shape)`, write a loop that on each step moves the current signal a small fraction `alpha` of the way toward `x_clean`, recording the `mse` against `x_clean` after every step. Run it for 50 steps with `alpha = 0.1` and print the error at the start, after one step, midway, and at the end. Then run it for only 5 steps and print the final error, so the effect of taking many small steps is visible.

```
# TODO 4: From pure noise, denoise toward x_clean in small steps of alpha=0.1, recording the mse
```

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Task 3 : Show Why Positional Information Is Needed

Write this cell from scratch. First show that bare attention is **order-blind**: compute the attention output for `X` and for the reordered sentence `X[perm]`, and check that reordering the words only reorders the outputs (the output for `X[perm]` equals the output for `X` with its rows reordered by `perm`). Then add positional information by using `X + P`, and repeat the same check on `X + P` versus `X[perm] + P` (the same words but now sitting at different positions). Print the boolean result of both checks so the contrast is visible.

```
# TODO 3: Print two booleans. First: bare attention is order-blind (reordering words only
#         reorders the outputs). Second: after adding P, that no longer holds (positions matter).
```

Task 4 : Denoise a Signal Step by Step

Write this cell from scratch. Model diffusion as cleaning up noise over many small steps toward the prompt's target. Starting from pure noise `rng.normal(size=x_clean.shape)`, write a loop that on each step moves the current signal a small fraction `alpha` of the way toward `x_clean`, recording the `mse` against `x_clean` after every step. Run it for 50 steps with `alpha = 0.1` and print the error at the start, after one step, midway, and at the end. Then run it for only 5 steps and print the final error, so the effect of taking many small steps is visible.

```
# TODO 4: From pure noise, denoise toward x_clean in small steps of alpha=0.1, recording the mse
```

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Task 4 : Denoise a Signal Step by Step

Write this cell from scratch. Model diffusion as cleaning up noise over many small steps toward the prompt's target. Starting from pure noise `rng.normal(size=x_clean.shape)`, write a loop that on each step moves the current signal a small fraction `alpha` of the way toward `x_clean` recording the `mse` against `x_clean` after every step. Run it for 50 steps with `alpha = 0.1` and print the error at the start, after one step, midway, and at the end. Then run it for only 5 steps and print the final error, so the effect of taking many small steps is visible.

```
# TODO 4: From pure noise, denoise toward x_clean in small steps of alpha=0.1, recording the mse
#         each step. Print the mse at start, after 1 step, midway, and at the end (50 steps),
#         and the final mse when only 5 steps are taken.
```

Task 5 : Explain What You Built (Text cell)

Add a **Text cell**. In your own words (4 to 6 sentences), explain all three of the following:

- What the attention weights represent and why each row **sums to one**, and what your Task 2 output shows about whether the most-attended word is always the nearest one (relevance versus distance).
- Why **bare attention is order-blind** and what **positional information** restores, referring to your two booleans from Task 3.

[Previous](#)

|| kalvium.community is sharing your screen. [Stop sharing](#) [Hide](#)

[Next](#)

Kalvium

Ask Gemini

kalvium.community/assessment/intro_genai4_v2_mpa3/attempt/c4435664-f366-4e37-bfc7-d2e41c97ba4a

School

 Foundations of Generative AI | MPA - 3 | 2026 - 27

56m left

Finish Assessment

Your camera feed, audio and screen are being proctored.



Instructions

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Task 5 : Explain What You Built (Text cell)

Add a **Text cell**. In your own words (4 to 6 sentences), explain all three of the following:

- What the attention weights represent and why each row **sums to one**, and what your Task 2 output shows about whether the most-attended word is always the nearest one (relevance versus distance).
- Why **bare attention is order-blind** and what **positional information** restores, referring to your two booleans from Task 3.
- How your Task 4 loop mirrors **diffusion**: starting from noise and removing a little each step toward the target, and what taking 5 steps instead of 50 did to the final error.

Submission Guidelines

- Ensure **every cell has been run** and all outputs and your Text cell are visible.
- In Colab: **File, then Print, then Destination Save as PDF, then Save.**
- Upload the single **PDF** here.

File naming:

Previous

kalvium.community is sharing your screen.


Stop sharing

Hide

Next

33°C Sunny

Search



ENG IN 12:05 09-10-2026

Kalvium

Ask Gemini

kalvium.community/assessment/intro_genai4_v2_mpa3/attempt/c4435664-f366-4e37-bfc7-d2e41c97ba4a

School

 Foundations of Generative AI | MPA - 3 | 2026 - 27

56m left

Finish Assessment

Your camera feed, audio and screen are being proctored.



Instructions

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Before You Submit : Checklist

- ☐ Solved entirely in **Google Colab**, in order
- ☐ Cell A ran and printed the tokens and the shapes
- ☐ Task 1: `softmax` and `attention` built from scratch; weight grid and row sums printed (rows sum to 1)
- ☐ Task 2: the most-attended word per word printed, and the row-sum confirmation
- ☐ Task 3: two booleans printed (bare attention order-blind True, with positions False)
- ☐ Task 4: mse printed at start, one step, midway, and end for 50 steps, and the final mse for 5 steps
- ☐ Task 5: a **Text cell** on attention weights, positional information, and diffusion
- ☐ Exported as **PDF** and named correctly

Upload your answer.

+ Add a file

Allowed file formats: pdf

Previous

kalvium.community is sharing your screen.

Stop sharing

[Hide](#)

Next

33°C Sunny

Search



ENG IN 12:05 09-10-2026

Assignment Report



Your report is here

 Score 10 / 10	 Accuracy 100%	 Time Spent 12s / 60m	 Attempted 1 / 1
--	--	---	---

Attempted on 08-Oct-2025 04:48 PM • Re-evaluation unavailable

Question 1 of 1

+10 0 0 10/10

Problem Statement

Real software is built by teams sharing one codebase. In this challenge you and a partner will collaborate on each other's repositories exactly the way a professional engineering team does: you will grant access, contribute a change through a pull request, review each other's work, and merge only after a review.

You must complete **both roles**: the **contributor** (you change your partner's repo) and the **reviewer** (you review the change your partner makes on your repo).

Before You Start

You need:

Assignment Report

Also answer, in 3 to 5 lines:

Why do protected branches and required reviews make a team safer than letting everyone push directly to main?

Uploaded File(s)



Team_contribution_submission.pdf
0.00 MB



What you did well?

All required tasks are completed with clear evidence in the screenshots. Each screenshot corresponds directly to the rubric criteria, demonstrating a thorough understanding of the collaboration process.

What you could do better?

No improvements needed; all aspects are well executed.



Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 1 of 21

What distinguishes a sequence from an unordered set?

- ☐ A sequence can only contain numbers, while a set can contain words
- ☐ A set is always larger than a sequence
- ☐ A sequence cannot contain repeated elements, but a set can
- ☐ In a sequence the order of elements is part of the information; in a set it is not

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 2 of 21

Which statement best captures why order information cannot simply be dropped?

- ☐ Order only matters for very long sentences
- ☐ Meaning in language depends on order, so losing order loses part of the meaning
- ☐ Order is decoration and does not affect meaning
- ☐ Order can always be guessed from word length

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 3 of 21

With positional encoding switched on, what becomes true of the same word appearing at two different positions?

- ☐ The two occurrences merge into a single word
- ☐ The two occurrences are deleted to save space
- ☐ The two occurrences are no longer identical to the model, because each carries its position
- ☐ The two occurrences swap their meanings

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

What is the main thing this unit says attention does NOT yet explain?

- ☐ That word order can carry meaning in a sentence
- ☐ That a pronoun can refer to an earlier noun
- ☐ How the relevance scores are actually computed from the words themselves
- ☐ That some words matter more than others for a given word

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

A model reads the money in the bank earned interest. Where should the contextual embedding of bank sit?

- ☐ Near words about rivers, such as water and shore
- ☐ Near words about finance, such as money and interest
- ☐ Exactly halfway between finance and river words at all times
- ☐ Nowhere in particular, since context is ignored

Clear Response

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Why can a strictly sequential model not process all elements in parallel?

- ☐ Because parallel processing would use too much electricity
- ☐ Because each step needs the summary produced by the previous step, so steps must run in order
- ☐ Because the elements are stored in alphabetical order
- ☐ Because it always finishes before the sequence is complete

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Why does distance between two related words weaken the link in a simple model?

- ☐ Because words far apart are always unrelated in meaning
- ☐ Because the fixed memory is re-compressed at every step, so far-apart words are harder to connect
- ☐ Because the model deletes any word it has already read
- ☐ Because longer distances make each word take up more memory

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 8 of 21

At a conceptual level, what does attention do for each word in a sentence?

- ☐ It weighs how relevant every other word is and blends them into a new representation of that word
- ☐ It removes every word except the most recent one
- ☐ It sorts the words into alphabetical order before reading
- ☐ It replaces the word with a fixed dictionary definition

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 9 of 21

In the sentences we sat on the grassy bank of the river and she saved money at the bank, what should differ about the embedding of bank?

- ☐ Its vector should differ, because the surrounding words point to different meanings
- ☐ Its spelling should change to mark the different meaning
- ☐ Nothing should differ, since it is the same word
- ☐ Only the length of the vector should change

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

- | | | |
|----|----|----|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | 9 |
| 10 | 11 | 12 |
| 13 | 14 | 15 |
| 16 | 17 | 18 |

Question 10 of 21

What problem does positional encoding solve for a model that reads all words at once?

- ☐ It makes the words shorter so they process faster
- ☐ It removes the need for the model to read the words
- ☐ It guarantees every sentence has the same meaning
- ☐ It restores the word order that the model would otherwise lose

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Your camera feed, audio and screen are being proctored.



Instructions

- | | | |
|----|----|----|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | 9 |
| 10 | 11 | 12 |
| 13 | 14 | 15 |
| 16 | 17 | 18 |

Question 11 of 21

What does it mean that bare self-attention is order-blind?

- ☐ It can only read words written in capital letters
- ☐ Reordering the words gives the same result, so it treats the sentence as an unordered bag of words
- ☐ It ignores the meaning of every word entirely
- ☐ It reads the sentence backwards on purpose

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Kalvium

Ask Gemini

kalvium.community/assessment/intro_genai4_v2_mpa3/attempt/c4435664-f366-4e37-bfc7-d2e41c97ba4a

School

 Foundations of Generative AI | MPA - 3 | 2026 - 27

 58m left

Finish Assessment

 Your camera feed, audio and screen are being proctored.



Instructions

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Question 12 of 21

In the cat that chased the dog was tired, who was tired?

☐ The dog, because it is the noun nearest to was tired

☐ Both the cat and the dog equally

☐ It cannot be determined from the sentence

☐ The cat, which is the subject introduced before the clause about the dog

Previous

kalvium.community is sharing your screen.


Stop sharing

[Hide](#)

Next

33°C Sunny

Search



ENG IN

12:03 09-10-2026



Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 13 of 21

What stays the same across the two senses of a word like light in different sentences?

- ☐ The vector the model produces for it
- ☐ The meaning the model assigns to it
- ☐ The neighbouring words that surround it
- ☐ The spelling of the word on the page

Clear Response

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 14 of 21

Which earlier idea does the contextual embedding build directly on?

- ☐ Sorting the words of a sentence into alphabetical order
- ☐ Counting how many characters each word contains
- ☐ Self-attention, where every word looks at the others in the same sentence
- ☐ Translating the sentence into another language first

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next





Your camera feed, audio and screen are being proctored.



Instructions

- | | | |
|----|----|----|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | 9 |
| 10 | 11 | 12 |
| 13 | 14 | 15 |
| 16 | 17 | 18 |

Question 15 of 21

What is the main limitation that contextual embeddings overcome compared with static ones?

- ☐ Static embeddings take up too much disk space to store
- ☐ Static embeddings cannot tell apart the different meanings of the same word
- ☐ Static embeddings can only represent numbers, not words
- ☐ Static embeddings change too often to be useful

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Kalvium

Ask Gemini

kalvium.community/assessment/intro_genai4_v2_mpa3/attempt/c4435664-f366-4e37-bfc7-d2e41c97ba4a

School

 Foundations of Generative AI | MPA - 3 | 2026 - 27

58m left

Finish Assessment

 Your camera feed, audio and screen are being proctored.



Instructions

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Question 16 of 21

In the engineer fixed the servers because they were failing, what does they refer to?

☐ The servers, a noun that appeared several words earlier

☐ The engineer, because that noun comes first in the sentence

☐ A new group of people not named in the sentence

☐ Nothing, because they never refers to a specific noun

Previous

kalvium.community is sharing your screen.


Stop sharing

[Hide](#)

Next

33°C Sunny

Search



ENG IN

12:04 09-10-2026



Your camera feed, audio and screen are being proctored.



Instructions

1	2	3
4	5	6
7	8	9
10	11	12
13	14	15
16	17	18

Question 17 of 21

What is a sequence, as the term is used in this unit?

- ☐ A collection of elements arranged in a definite order, where the order carries meaning
- ☐ An unordered set of elements where arrangement makes no difference
- ☐ A single number that summarises a whole piece of data
- ☐ A grid of values with rows and columns but no order

Previous

kalvium.community is sharing your screen.

Stop sharing

Hide

Next



Kalvium

Ask Gemini

kalvium.community/assessment/intro_genai4_v2_mpa3/attempt/c4435664-f366-4e37-bfc7-d2e41c97ba4a

School

 Foundations of Generative AI | MPA - 3 | 2026 - 27

58m left

Finish Assessment

 Your camera feed, audio and screen are being proctored.



Instructions

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

Question 18 of 21

Why does reading all words at once create a problem for order?

☐ Taken together with no sequence, the words arrive as a set with no built-in order

☐ Reading all at once makes the model forget the words themselves

☐ Reading all at once changes the spelling of the words

☐ Reading all at once always reverses the sentence

Previous

kalvium.community is sharing your screen.

Stop sharing

[Hide](#)

Next

33°C Sunny

Search



ENG IN

12:04 09-10-2026